

Purple Team — Weekender: Pittsburgh

Plan-O-Matic

Executive Summary

Goal. Build a Python-only command-line tool that plans an optimal Pittsburgh weekend by scoring and ranking events against forecasted weather and nearby food options. The output is a ranked itinerary plus clear visuals (charts and an optional map) that look polished without heavy engineering.

Why it matters. People waste time juggling tabs for events, weather, and restaurants. We unify them into one repeatable data pipeline and a simple decision algorithm.

Data plan (≥ 3 sources, ≥ 1 scraped). - **Weather:** Open-Meteo forecast (hourly/daily; no key) - **Events:** Eventbrite Events API (with token) **and** a scraped local calendar page (e.g., VisitPittsburgh) - **Food nearby:** Yelp Fusion Business Search API (to count/rank well-rated spots close to each event)

Deliverable. A single runnable Python entrypoint (e.g., `weekender.py`) that prompts for date/time window, budget, and start location, then prints and saves: (1) a ranked event list with component scores; (2) charts as PNGs; (3) optional HTML map if we include a Folium view.

Project Objectives

1. **Aggregate** events for a chosen weekend (Fri–Sun) from at least two distinct sources (API + one scraped page).
 2. **Enrich** each event with weather features matching its time window and with nearby dining density/quality.
 3. **Score** each event via a transparent composite formula (weather fit, food density/quality, travel distance, budget fit).
 4. **Output** a sorted itinerary and simple visuals, with a cache option to avoid slow/fragile network calls.
 5. **Document** a clean setup/run path (README), and provide a first-draft pitch deck aligned to the class template.
-

Users & Use Cases

- **Students / residents** deciding between events when the weather is iffy.
 - **Visitors** who want a quick “best bet” list for a specific date with tasty food nearby.
 - **Policy/Events teams** (stretch): sense-check whether outdoor programming is robust to weather.
-

Scope & Non-Goals

In scope - Weekends in the next ~2 weeks (configurable) - Simple heuristic classification of indoor/outdoor events - Straight-line travel time proxy (haversine + default walking/driving speed) - CSV/JSON cache to support reproducible runs

Out of scope (for v1) - True routing/traffic times, ticket purchasing, calendars integration - Advanced recommender models or user history

Data Sources (Planned)

1. **Open-Meteo API** — hourly precipitation probability, temperature, wind. Used to derive a per-event **weather_fit** score.
2. **Eventbrite API** — event lists, times, venues (lat/lon where available). Used for canonical event attributes.
3. **Scraped local events page** — VisitPittsburgh (or Downtown Partnership/Cultural Trust calendar) to meet the scraping requirement and to supplement the API with hyperlocal events.
4. **Yelp Fusion API** — dining options near event venue; derive **food_score** from counts and ratings within a radius.

Data freshness. Weather is fetched for the specific weekend (or read from cache if `--use-cache`). Events and Yelp data are cached per run.

Robustness. If scraping fails, the program still runs using Eventbrite + Yelp. If an API is rate-limited, we back off and use cached JSON.

System Architecture (Python-only)

```
weekender/  
  weekender.py          # SINGLE entrypoint (CLI). Imports modules  
below.  
  config.py             # API tokens, defaults, constants (reads  
from .env if present)  
  etl/  
    fetch_eventbrite.py  # API pull, with pagination & basic  
normalization  
    scrape_visitpgh.py   # Requests + BeautifulSoup scrape → rows [title,  
dt, venue, url]  
    fetch_weather.py     # Open-Meteo hourly forecast retrieval  
    fetch_yelp.py        # Yelp search near (lat, lon) for eateries  
    cache.py             # read/write JSON & CSV; simple TTL logic  
features/  
  classify_indoor_outdoor.py # keyword rules ("park", "trail", "museum",
```

```

etc.)
    scoring.py                # composite score; helpers for normalization
    geo.py                    # haversine distance; coarse travel-time proxy
viz/
    charts.py                 # matplotlib figures saved to /out/*.png
    mapview.py                # optional Folium HTML map (stretch goal)
data/
    raw/                     # cached API/HTML responses
    processed/                # tidy CSVs
out/                           # report CSV + PNG charts + optional HTML map
README.md
requirements.txt

```

Stack & Libraries

- **Core:** `python>=3.11`, `requests`, `pandas`, `beautifulsoup4`, `matplotlib`
- **Optional:** `folium` (HTML map), `python-dotenv` (local tokens), `tqdm` (progress bars)
- **Testing/quality (light):** `pytest`, `black`, `ruff` (optional but recommended)

Why this stack? It's standard, lightweight, and satisfies the course emphasis (scraping + APIs + pandas + matplotlib) while keeping installation minimal.

CLI Interface (User Flow)

Run:

```

python weekender.py --start 2025-09-12 --end 2025-09-14
--home "5000 Forbes Ave, Pittsburgh, PA"
--budget 60 --style outdoorsy --fresh yes

```

Prompts (if flags omitted): date range, budget (\$), preferred style (`indoorsy|outdoorsy|mixed`).

Outputs: - `out/ranked_events.csv` with columns: `rank,score,weather_fit,food_score,travel_score,budget_score,title,start,end,venue,lat,lon,source,u`
 - `out/top10_bar.png`, `out/timeline.png` (and `out/map.html` if Folium is included) - Console printout of Top 5 with component breakdowns

Feature Engineering & Scoring

1) Indoor/Outdoor classification (rule-based)

- **Outdoor hints:** title/venue contains {park, trail, riverfront, outdoors, plaza, market (outdoor), concert (outdoor)}
- **Indoor hints:** {museum, gallery, theater, library, hall, arena, rink, studio}
- Default: `indoors` when unknown.

2) Weather features (from hourly forecast)

For each event time window, aggregate: - `p_precip` = max precipitation probability in window - `t_mean` = average temperature in window (°F) - `wind_max` = max wind speed in window

3) Weather-fit score

Normalize each component to 0-1 and compute:

```
weather_fit = 1
if outdoor:
    weather_fit -= 0.7 * clamp01((p_precip - 0.15)/0.35) # strong rain penalty
    # above 15%
    weather_fit -= 0.2 * clamp01((abs(t_mean-72))/25)      # discomfort away from
    # ~72°F
    weather_fit -= 0.1 * clamp01((wind_max-10)/20)        # breezy penalty
else: # indoor
    weather_fit -= 0.3 * clamp01((p_precip - 0.30)/0.50) # smaller rain penalty
    weather_fit -= 0.2 * clamp01((abs(t_mean-72))/35)
weather_fit = clamp01(weather_fit)
```

4) Food score (Yelp)

Within radius `R=600m` of venue: count businesses with `rating≥4.2` and `review_count≥50`.

```
food_score = clamp01( 0.7 * min(top_spots, 8)/8 + 0.3 * avg_rating/5 )
```

5) Travel score (coarse)

Haversine distance from `home` to venue. Assume 5 km/h walking or 25 km/h rideshare default.

```
travel_minutes = (distance_km / 25) * 60
travel_score = clamp01(1 - travel_minutes/30) # 30+ min decays to ~0
```

6) Budget score (simple)

If event has price info: full points if `price <= budget`, else linear decay until `2*budget`. Unknown price → neutral (0.6) so unknowns aren't unfairly punished.

7) Composite score (weights)

```
score = 0.45*weather_fit + 0.25*food_score + 0.15*travel_score +  
0.15*budget_score
```

Weights are easy to tweak in `config.py`.

Data Model (Tidy Tables)

events.csv - `event_id, source (eventbrite|visitpgh), title, description, category, start, end, venue_name, lat, lon, price_low, price_high, url, indoor_outdoor`

weather_hourly.csv - `ts, lat, lon, temp_f, precip_prob, wind_mph`

yelp_stats.csv - `event_id, radius_m, top_spots, avg_rating, avg_review_count`

ranked_events.csv (final) - `rank, score, weather_fit, food_score, travel_score, budget_score, <event fields...>`

Roles & Workstreams (Team of 4)

Purple Team Roster - **Noah** — Project Lead & Viz/Packaging - **Teammate B** — Events Ingestion (API + Scrape) - **Teammate C** — Weather & Scoring - **Teammate D** — Yelp & Geospatial

Swap names as you wish. Responsibilities below are designed to interlock with minimal merge pain.

Role 1 — Events Ingestion (API + Scrape)

Deliverables - `etl/fetch_eventbrite.py` with: token auth, city/lat-lon filters, pagination, venue lat/lon resolution - `etl/scrape_visitpgh.py` with: HTTP GET, polite headers, table/list parsing via BeautifulSoup, date normalization - `data/processed/events.csv` (deduped across sources)

Key details - Normalize time to local TZ (America/New_York). Ensure ISO-8601 strings. - Use a **source** column to track provenance. - Implement simple cleaning: strip whitespace, drop events missing title/start. - Caching: write raw responses to `data/raw/eventbrite_*.json`, `data/raw/visitpgh_*.html`.

Tests - Parsing unit tests on static HTML sample (saved under `tests/samples/visitpgh.html`).

Role 2 — Weather & Scoring

Deliverables - `etl/fetch_weather.py` that pulls hourly weather for each unique venue (lat, lon) and the event time window - `features/scoring.py` implementing the formula above + normalization helpers - `data/processed/weather_hourly.csv`

Key details - Map event time windows to relevant hourly rows (left-join by hour). - Ensure temperature is Fahrenheit and wind in mph (convert if needed). - Provide a `--fresh` flag to bypass cache.

Tests - Deterministic scoring for crafted inputs (e.g., heavy rain → low score).

Role 3 — Yelp & Geospatial

Deliverables - `etl/fetch_yelp.py` to query food near each venue, returning `top_spots`, `avg_rating`, etc. - `features/geo.py` providing haversine + travel minutes. - `data/processed/yelp_stats.csv`

Key details - Respect daily request caps; sleep between pages. - If venue lat/lon missing, geocode fallback is **out of scope**; skip with a warning.

Tests - Deterministic extraction of `top_spots` from a saved JSON fixture.

Role 4 — Viz/Packaging & CLI

Deliverables - `weekender.py` (single entrypoint) invoking modules and writing outputs - `viz/charts.py` producing: `top10_bar.png`, `timeline.png` - Optional `viz/mapview.py` producing `map.html` (Folium) - `README.md` with install/run steps and screenshots of outputs

Key details - Argument parsing with `argparse` and sensible defaults. - Clear console output (Top 5 with component scores). - Save artifacts to `/out` with timestamped filenames.

Tests - Smoke test: run CLI on fixture data with `--use-cache` to produce outputs without network.

Git & Collaboration

- **Repo name:** `purple-weekender`
 - **Branching:** feature branches per role (`feat/events`, `feat/weather`, `feat/yelp`, `feat/viz`)
 - **PRs:** small, reviewer from adjacent role; squash-merge
 - **Code style:** apply `black` and `ruff` (pre-commit optional)
 - **Issues/Milestones:** use GitHub Projects board (To Do / In Progress / Review / Done)
-

Schedule & Milestones

By Thursday (Draft deck due) - Events ingestion MVP from **one** source (Eventbrite **or** scraper) with cache - Weather fetch for a single venue and date, working end-to-end - CLI that prints Top 5 using stubbed/default scores - One chart saved as PNG (e.g., `top10_bar.png` on dummy data) - First-draft slides (Vision → Problem → Solution → Data Sources → Demo preview)

Week 2 - Add second/third data sources (scraped site + Yelp) - Implement full scoring and two charts + optional map - README v1 with screenshots

Week 3 (polish) - Robust caching and error handling, final weights tuning, UX polish - Final slide deck + short demo video

Risks & Mitigations

- **Scrape breaks** → Pin a sample HTML file and support cached parse.
 - **API rate limits** → Cache responses; exponential backoff; document how to add tokens.
 - **Missing venue coords** → Skip event with warning; plenty of others remain.
 - **Weather mismatch** → Always log the chosen hourly rows per event for auditability.
-

Ethics, Terms & Privacy

- Respect `robots.txt` and site terms. Identify user-agent and add a polite delay for scraping.
 - Only collect public event info; no PII. Keep tokens in local `.env` (not committed).
-

Installation & Running (Draft)

requirements.txt (initial)

```
requests
pandas
beautifulsoup4
matplotlib
folium
python-dotenv
```

Setup

```
python -m venv .venv
source .venv/bin/activate # Windows: .venv\Scripts\activate
pip install -r requirements.txt
cp .env.example .env # add EVENTBRITE_TOKEN, YELP_API_KEY
python weekender.py --start 2025-09-12 --end 2025-09-14 --home "5000 Forbes Ave, Pittsburgh, PA"
```

`.env.example`

```
EVENTBRITE_TOKEN=xxxxxxxxxxxxxxxxxx
YELP_API_KEY=xxxxxxxxxxxxxxxxxx
```

Minimal Code Skeleton (illustrative)

```
# weekender.py
from etl.fetch_eventbrite import get_events
from etl.scrape_visitpgh import scrape_events
from etl.fetch_weather import get_weather
from etl.fetch_yelp import get_yelp_stats
from features.scoring import score_events
from viz.charts import save_top10_bar, save_timeline

# parse args → call pipeline steps → write ranked_events.csv and charts
```

Slide Deck Mapping (for the Draft)

- **Vision:** “Plan the perfect Pittsburgh weekend in one go.”
- **Problem:** Tabs + gut feel → suboptimal choices.
- **Solution:** Unified Python tool, transparent scoring, pretty outputs.
- **Data:** list sources; call out which is scraped; show a small table sample.
- **Demo Preview:** screenshot of `ranked_events.csv` head + `top10_bar.png`.
- **Team:** Purple Team roles + who is spokesperson and Canvas submitter.

Definition of Done (DoD)

- One command runs end-to-end with either fresh or cached data
- `ranked_events.csv` produced with all four component scores
- At least two distinct charts saved to `/out`

- README with installation/run + screenshots; tokens handled via `.env`
 - Draft and final pitch decks complete and consistent with the tool's outputs
-

Stretch Ideas (if time allows)

- Lightweight clustering of events (time/price/distance) to suggest combos
 - Simple “day planner” that avoids overlapping start times
 - Add a preference profile (indoors vs outdoors weight) persisted to a small JSON file
-

Purple Team FTW. Simple pipeline, fancy outputs, clean code. Let's ship it. 💜